

PARALLEL AND DISTRIBUTED COMPUTING

LAB 2

S Shyam Sundaram

19BCE1560

18th August, 2021

Dr SK Ayesha

L43+L44

AIM

Compare the time taken when matrix addition is serialized and parallelized using Amdahl's Law to calculate the speedup.

CODE

```
#include <stdio.h>
#include "omp.h"
#include<time.h>

#define ROWS 2500
#define COLS 250

int main()
{
    int a[ROWS][COLS],b[ROWS][COLS],c[ROWS][COLS];

    for(int i=0;i<ROWS;++i)
    for(int j=0;j<COLS;++j)
    {
        a[i][j]=i*10+j;
        b[i][j]=j*10+i;
    }

    float start=omp_get_wtime();

    #pragma omp parallel for shared(a,b,c) //reduction(+: c)
    for(int i=0;i<ROWS;++i)
    for(int j=0;j<COLS;++j)
    {
        c[i][j]=a[i][j]+b[i][j];
    }
    float end=omp_get_wtime();
    float exec=end-start;

    printf("Time taken %f\n",exec);
    return 0;
}
```

EXECUTION

Bash commands:

```
export OMP_NUM_THREADS=2
gcc -fopenmp prog.c
./a.out
```

Note: For Sequential, OMP_NUM_THREADS=1

RESULTS

We see that serial execution takes a lot longer than most of the parallelized implementations. The average execution time reduces as the number of threads increase from 2 to 6. But, from thread count 8, we see a general increase in the time taken to execute matrix addition. The observations are recorded in the table below.

```
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ export OMP_NUM_THREADS=1
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ gcc -fopenmp matadd.c
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.009766
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.007812
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.007812
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.007812
```

```
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ export OMP_NUM_THREADS=1
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ gcc -fopenmp matadd.c
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.009766
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.007812
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.007812
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab2$ ./a.out
Time taken 0.007812
```

TABLE

Thread Count	Execution Time 1	Execution Time 2	Execution Time 3	Execution Time 4	Average Execution Time
1	0.009766	0.007812	0.011719	0.001953	0.007812
2	0.003906	0.003906	0.003906	0.004883	0.004150
4	0.001953	0.003906	0.001953	0.003906	0.002930
6	0.001953	0.001953	0.003906	0.001953	0.002443
8	0.001953	0.001953	0.001953	0.015625	0.005371
16	0.003906	0.003906	0.003906	0.003906	0.003906
32	0.003906	0.005859	0.005859	0.005859	0.005371
64	0.005859	0.019531	0.007812	0.007812	0.010226

Thread Count	PF	Speed Up
2	0.937532	1.88241
4	0.833248	2.66621
6	0.824731	3.19771
8	0.357106	1.45448
16	0.533333	1.99999
32	0.322548	1.45447
64	-0.313917	0.76394

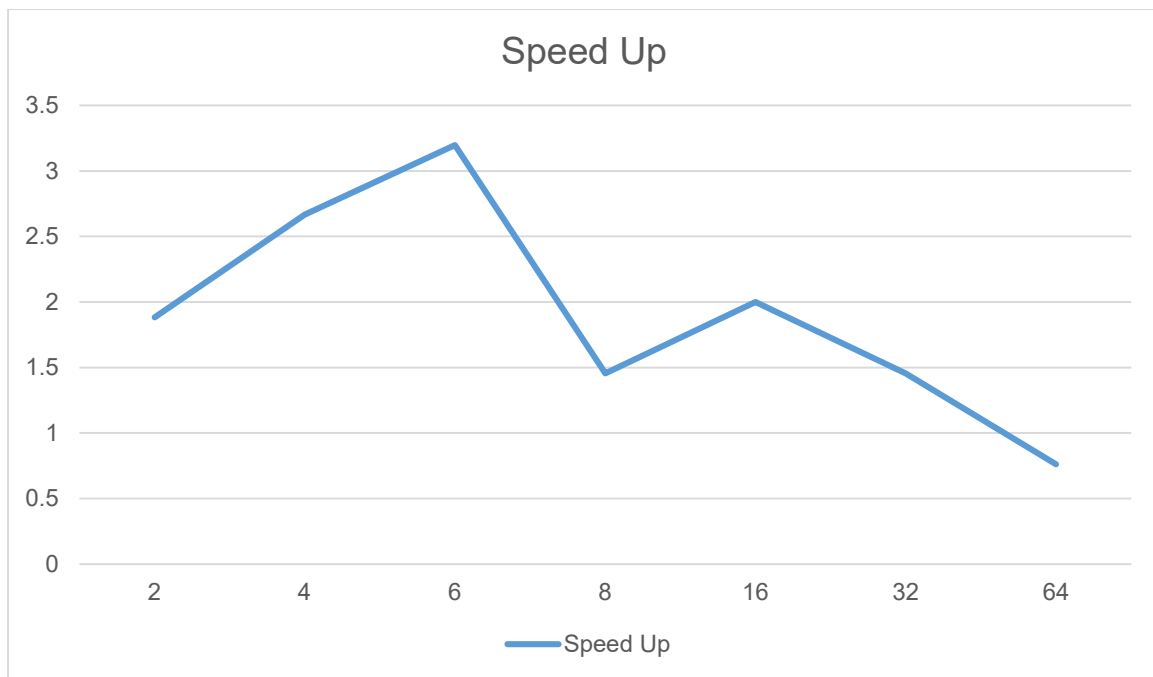
We see that the optimal number of threads is 6, so $p=6$.

$T_p=0.002443$ and $T_1=0.007812$. Thus, $PF = (1 - \frac{T_p}{T_1}) / (1 - \frac{1}{p}) = 0.824731$

From Amdahl's law, $speedup = 1 / (s + \frac{p}{n})$, where $p = PF$, $s=1-p$ and n =optimal num of threads which here is 6. Which gives us a speedup = 3.197706.

GRAPH





CONCLUSION

We have thus calculated the speed up to be nearly 3.2 when parallelizing the problem of matrix addition for large matrices.