

PARALLEL AND DISTRIBUTED COMPUTING LAB

REPORT

NAME: S Shyam Sundaram

REG NO: 19BCE1560

PROGRAMMING ENVIRONMENT: OpenMP

PROBLEM: Vector and Matrix Addition

DATE: 25th August, 2021

HARDWARE CONFIGURATION:

CPU NAME	:	Intel core i5 – 1035G1 @ 1.00 Ghz
Number of Sockets:	:	1
Cores per Socket	:	4
Threads per core	:	1
L1 Cache size	:	320KB
L2 Cache size	:	2MB
L3 Cache size (Shared):	:	6MB
RAM	:	8 GB

VECTOR ADDITION

CODE

```
#include <stdio.h>
#include "omp.h"
#include <time.h>

#define N 600000

int main()
{
    float a[N],b[N],c[N];
    int i;
    float start,end,exec;
    printf("Name: Shyam Sundaram\nReg num: 19BCE1560\nPDC Lab:\n\n");

    for(i=0;i<N;++i)
    {

        a[i]=(i+1)*1.0;
        b[i]=(i+1)*2.0;
    }

    int thread[]={1,2,4,8,16,32,64,128,256,512};
```

```

float serial;
for(int t=0;t<10;++t)
{
    omp_set_num_threads(thread[t]);
    start=omp_get_wtime();
    #pragma omp parallel default(none), private(i,m), shared(a,b,c)
    {
        #pragma omp for
        for(i=0;i<N;++i)
        {
            for(int j=0;j<1000;++j) //m is 1000 here to increase workload
            c[i]=a[i]+b[i];
        }
    }
    end=omp_get_wtime();
    exec=end-start;
    if(t==0) serial=exec;
    printf("Thread count: %d Time taken is: %f ",thread[t],exec);
    float pf=(1-(exec/serial))/(1-(1/thread[t]));
    printf(" PF = %f ",pf);
    float s=1-pf;
    float speedup=1/(s+(pf/thread[t]));
    printf(" Speedup = %f\n",speedup);
}
return 0;
}

```

COMPILE AND EXECUTION

gcc -fopenmp three.c
./a.out

OBSERVATIONS

NUMBER OF THREADS	EXECUTION TIME	SPEED-UP
1	6.472656	1
2	0.729736	1.797363
4	0.367920	3.417265
8	0.351318	5.797349
16	0.358398	8.740459
32	0.353760	11.876975
64	0.357178	14.296876
128	0.357422	15.974100
256	0.363037	16.729416
512	0.363770	17.228210

ASSUMPTION


Following extra for loop is added to increase the number of operations in the parallel region to visualize the effect of multi-threading in vector addition.


```
for(int j=0;j<m;j++)  
  
    c[i] = a[i] + b[i];
```

SCREENSHOT

```
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab3$ gcc -fopenmp three.c  
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab3$ ./a.out  
Name: Shyam Sundaram  
Reg num: 19BCE1560  
PDC Lab:  
  
Thread count: 1 Time taken is: 6.472656 PF = -nan Speedup = -nan  
Thread count: 2 Time taken is: 0.729736 PF = 0.887259 Speedup = 1.797363  
Thread count: 4 Time taken is: 0.367920 PF = 0.943158 Speedup = 3.417265  
Thread count: 8 Time taken is: 0.351318 PF = 0.945723 Speedup = 5.797349  
Thread count: 16 Time taken is: 0.358398 PF = 0.944629 Speedup = 8.740459  
Thread count: 32 Time taken is: 0.353760 PF = 0.945346 Speedup = 11.876975  
Thread count: 64 Time taken is: 0.357178 PF = 0.944817 Speedup = 14.296876  
Thread count: 128 Time taken is: 0.357422 PF = 0.944780 Speedup = 15.974100  
Thread count: 256 Time taken is: 0.363037 PF = 0.943912 Speedup = 16.729416  
Thread count: 512 Time taken is: 0.363770 PF = 0.943799 Speedup = 17.228210
```

PLOTS





INFERENCE

The addition of the extra for loop increased the workload. Thus, as a greater number of threads work on it, the lower the execution time is and higher the Speed-Up, but up to a certain point, after which it is near constant.

MATRIX ADDITION

CODE

```
#include <stdio.h>
#include "omp.h"
#include <time.h>

#define ROWS 2500
#define COLS 250

int main()
{
    float a[ROWS][COLS], b[ROWS][COLS], c[ROWS][COLS];
    printf("Name: Shyam Sundaram\nReg num: 19BCE1560\nPDC Lab:\n\n");

    for(int i=0; i<ROWS; ++i)
        for(int j=0; j<COLS; ++j)
        {
            a[i][j] = i*10+j;
            b[i][j] = j*10+i;
        }
    int thread[] = {1, 2, 4, 8, 16, 32, 64, 128, 256, 512};
    float serial;
    for(int t=0; t<10; ++t)
```

```

{
    omp_set_num_threads(thread[t]);
    float start=omp_get_wtime();

    #pragma omp parallel for shared(a,b,c) //reduction(+: c)
    for(int i=0;i<ROWS;++i)
    for(int j=0;j<COLS;++j)
    {
        for(int j=0;j<1000;++j)
            c[i][j]=a[i][j]+b[i][j];
    }
    float end=omp_get_wtime();
    float exec=end-start;
    if(t==0) serial=exec;

    printf("Thread count: %d Time taken is: %f",thread[t],exec);
    float pf=(1-(exec/serial))/(1-(1/thread[t]));
    printf(" PF = %f ",pf);
    float s=1-pf;
    float speedup=1/(s+(pf/thread[t]));
    printf(" Speedup = %f\n",speedup);
}
return 0;
}

```

COMPILATION AND EXECUTION

gcc -fopenmp matadd.c
./a.out

OBSERVATIONS

NUMBER OF THREADS	EXECUTION TIME	SPEED-UP
1	8.358765	1
2	2.782715	1.500477
4	0.533691	3.356988
8	0.492065	5.665413
16	0.508057	8.369429
32	0.503052	11.166713
64	0.512817	13.154904
128	0.507690	14.689577
256	0.515747	15.298335
512	0.520386	15.603592

ASSUMPTION


Following extra for loop is added to increase the number of operations in the parallel region to visualize the effect of multi-threading in vector addition.


```
for(int j=0;j<m;j++)  
  
    c[i] = a[i] + b[i];
```

SCREENSHOT

```
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab3$ gcc -fopenmp matadd.c  
shyam@shyam-Inspiron-14-5408:~/Academics/Labs/PDC/Lab3$ ./a.out  
Name: Shyam Sundaram  
Reg num: 19BCE1560  
PDC Lab:  
  
Thread count: 1 Time taken is: 8.358765 PF = -nan Speedup = -nan  
Thread count: 2 Time taken is: 2.782715 PF = 0.667090 Speedup = 1.500477  
Thread count: 4 Time taken is: 0.533691 PF = 0.936152 Speedup = 3.356988  
Thread count: 8 Time taken is: 0.492065 PF = 0.941132 Speedup = 5.665413  
Thread count: 16 Time taken is: 0.508057 PF = 0.939219 Speedup = 8.369429  
Thread count: 32 Time taken is: 0.503052 PF = 0.939817 Speedup = 11.166713  
Thread count: 64 Time taken is: 0.512817 PF = 0.938649 Speedup = 13.154904  
Thread count: 128 Time taken is: 0.507690 PF = 0.939263 Speedup = 14.689577  
Thread count: 256 Time taken is: 0.515747 PF = 0.938299 Speedup = 15.298335  
Thread count: 512 Time taken is: 0.520386 PF = 0.937744 Speedup = 15.603592
```

PLOTS





INFERENCE

The addition of the extra for loop increased the workload. Thus, as a greater number of threads work on it, the lower the execution time is and higher the Speed-Up, but up to a certain point, after which it is near constant.